

Untitled

Yue Lai

5/16/2020

import data

```
red = read_excel(path = "./data/wine.xlsx", sheet = "red") %>%
  janitor::clean_names()

## Warning in FUN(X[[i]], ...): strings not representable in native encoding will
## be translated to UTF-8

white = read_excel(path = "./data/wine.xlsx", sheet = "white") %>%
  janitor::clean_names()

wine = data.frame(rbind(red, white))

wine = wine %>%
  mutate(quality = as.factor(ifelse(quality > 5, "good", "bad")))
```


Divide the data into two part (training and test)


```
rowtrain = createDataPartition(y = wine$quality,
                                p = 2/3,
                                list = FALSE)

ctrl = trainControl(method = "repeatedcv",
                     repeats = 5,
                     summaryFunction = twoClassSummary,
                     classProbs = TRUE)
```

tree

```
## using caret
rpart.fit = train(quality~., wine,
                  subset = rowtrain,
                  method = "rpart",
                  tuneGrid = data.frame(cp = exp(seq(-6, -2, len = 20))),
                  trControl = ctrl,
                  metric = "ROC")
ggplot(rpart.fit, highlight = TRUE)
```






```
library(rpart)
tree1 = rpart(formula = quality ~ ., data = wine,
               subset = rowtrain,
               control = rpart.control(cp = 0))
cpTable = printcp(tree1)
```

```
##
## Classification tree:
## rpart(formula = quality ~ ., data = wine, subset = rowtrain,
##       control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] alcohol  cacid   chlorides density  facid   freesd  p_h
## [8] rsugar   sulphates totalsd vacid
##
## Root node error: 1590/4332 = 0.36704
##
## n= 4332
##
##      CP nsplit rel error  xerror   xstd
## 1 0.13710692    0  1.00000 1.00000 0.019952
## 2 0.01194969    2  0.72579 0.74340 0.018438
## 3 0.00723270    4  0.70189 0.71950 0.018249
## 4 0.00691824    6  0.68742 0.72013 0.018254
## 5 0.00628931    7  0.68050 0.72075 0.018259
```


## 6	0.00503145	8	0.67421	0.70943	0.018166
## 7	0.00440252	9	0.66918	0.70252	0.018108
## 8	0.00408805	11	0.66038	0.69748	0.018066
## 9	0.00356394	13	0.65220	0.68616	0.017968
## 10	0.00314465	29	0.56918	0.67862	0.017902
## 11	0.00272537	30	0.56604	0.67170	0.017841
## 12	0.00251572	35	0.55220	0.67107	0.017835
## 13	0.00235849	50	0.51258	0.66352	0.017767
## 14	0.00230608	54	0.50314	0.66352	0.017767
## 15	0.00220126	57	0.49623	0.66289	0.017762
## 16	0.00188679	61	0.48742	0.66730	0.017802
## 17	0.00157233	73	0.46352	0.67925	0.017908
## 18	0.00125786	89	0.43459	0.68050	0.017919
## 19	0.00104822	100	0.42075	0.68113	0.017925
## 20	0.00094340	103	0.41761	0.69308	0.018028
## 21	0.00083857	109	0.41195	0.69434	0.018039
## 22	0.00078616	112	0.40943	0.69623	0.018055
## 23	0.00075472	116	0.40629	0.69623	0.018055
## 24	0.00062893	126	0.39874	0.69937	0.018082
## 25	0.00041929	145	0.38428	0.69937	0.018082
## 26	0.00031447	148	0.38302	0.69874	0.018076
## 27	0.00020964	156	0.38050	0.70943	0.018166
## 28	0.00015723	159	0.37987	0.71509	0.018213
## 29	0.00000000	163	0.37925	0.71950	0.018249

```
plotcp(tree1)
```



```
minErr = which.min(cpTable[,4])

tree2 = prune(tree1, cp = cpTable[minErr, 1])
rpart.plot(tree2)
```



bagging and random forests

```
## using caret
rf.grid = expand.grid(mtry = 1:6,
#                       splitrule = "gini",
#                       min.node.size = 1:6)

#rpart.fit = train(quality~., wine,
#                  subset = rowtrain,
#                  method = "ranger",
#                  tuneGrid = rf.grid,
#                  trControl = ctrl,
#                  metric = "ROC")

#ggplot(rpart.fit, highlight = TRUE)
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.6.3
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ranger':
```

```
##
```

```
##      importance
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      combine
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

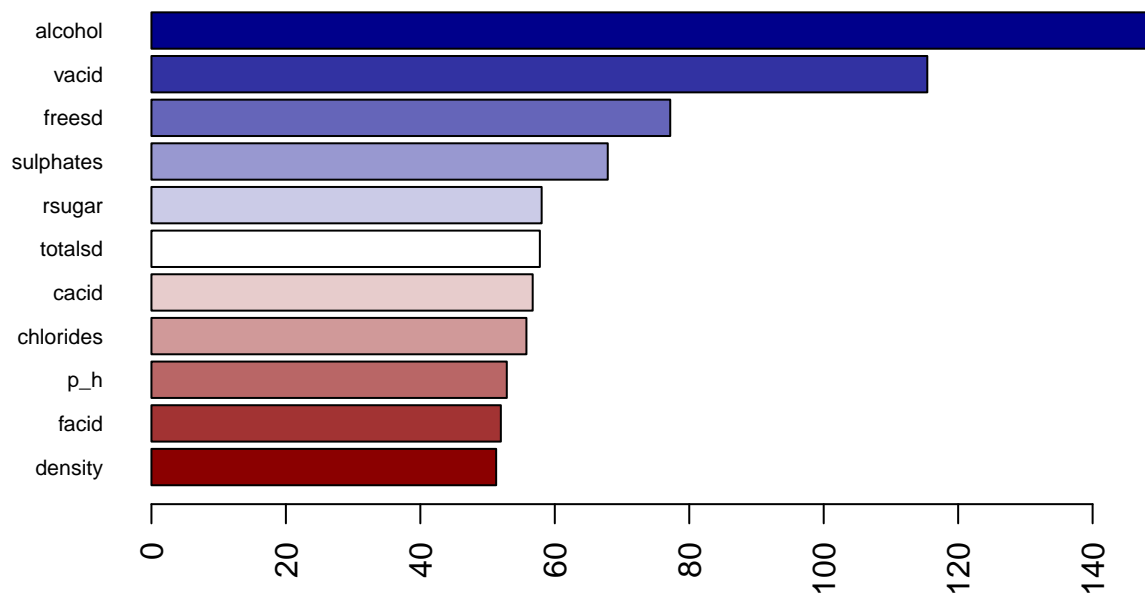
```
##      margin
```

```
# bagging
```

```
bagging = randomForest(quality~., wine[rowtrain,], mtry = 11)
```

```
bagging.per = ranger(quality~., wine[rowtrain,],  
                     mtry = 11,  
                     splitrule = "gini",  
                     min.node.size = 5,  
                     importance = "permutation",  
                     scale.permutation.importance = TRUE)
```


```
barplot(sort(ranger::importance(bagging.per), decreasing = FALSE),  
        las = 2, horiz = TRUE, cex.names = 0.7,  
        col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(11))
```



```
# random forest
rf = randomForest(quality~., wine[rowtrain,],mtry = 3)

rf.per = ranger(quality~., wine[rowtrain,],
                mtry = 3,
                splitrule = "gini",
                min.node.size = 5,
                importance = "permutation",
                scale.permutation.importance = TRUE)

barplot(sort(ranger::importance(rf.per), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("darkred", "white", "darkblue"))(11))
```




boosting

```
library(gbm)
```


```
## Warning: package 'gbm' was built under R version 3.6.3
```

```
## Loaded gbm 2.1.5
```

```
wine2 = wine
wine2$quality = as.numeric(wine2$quality == "good")
bst = gbm(quality~., wine2[rowtrain,],
  distribution = "adaboost",
  n.trees = 20000,
  interaction.depth = 3,
  shrinkage = 0.001,
  cv.folds = 10)
nt = gbm.perf(bst, method = "cv")
```

```
summary(bst, las = 2, cBars = 8, cex.names = 0.6)
```



```
##           var  rel.inf
## alcohol    alcohol 26.948308
## vacid      vacid  18.317838
## freesd     freesd  8.930802
## totalsd    totalsd 8.371572
## sulphates  sulphates 7.827669
## rsugar     rsugar  6.827096
## density    density  5.008816
## chlorides  chlorides 4.966148
## cacid      cacid   4.876538
## facid      facid   4.002520
## p_h        p_h    3.922693
```

boosting in train

```
## (long time to run )
#gbm.grid_2 <- expand.grid(n.trees = c(2000,3000,4000),
#                          interaction.depth = 1:6,
#                          shrinkage = c(0.001, 0.003,0.005),
#                          n.minobsinnode = 1)
#gbm.fit_2 <- train(quality~., wine[rowtrain,],
#                  method = "gbm",
#                  tuneGrid = gbm.grid_2,
#                  trControl = ctrl,
```

```

#           metric = "ROC",
#           verbose = FALSE)
#ggplot(gbm.fit_2, highlight = TRUE)

tree.pred = predict(tree2, newdata = wine[-rowtrain,], type = "prob")[,1]
bag.pred = predict(bagging, newdata = wine[-rowtrain,], type = "prob")[,1]
rf.pred = predict(rf, newdata = wine[-rowtrain,], type = "prob")[,1]

bst.pred = predict(bst, newdata = wine[-rowtrain,], type = "response")

## Using 20000 trees...

roc.tree = roc(wine$quality[-rowtrain], tree.pred)

## Setting levels: control = bad, case = good

## Setting direction: controls > cases

roc.bag = roc(wine$quality[-rowtrain], bag.pred)

## Setting levels: control = bad, case = good
## Setting direction: controls > cases

roc.rf = roc(wine$quality[-rowtrain], rf.pred)

## Setting levels: control = bad, case = good
## Setting direction: controls > cases

roc.bst = roc(wine$quality[-rowtrain], bst.pred)

## Setting levels: control = bad, case = good

## Setting direction: controls < cases


auc = c(roc.tree$auc[1], roc.bag$auc[1], roc.rf$auc[1], roc.bst$aur[1])

plot(roc.tree, legacy.axes = TRUE)
plot(roc.bag, col = 2, add = TRUE)
plot(roc.rf, col = 3, add = TRUE)
plot(roc.bst, col = 4, add = TRUE)

modelName = c("tree", "bagging", "rf", "boosting")

legend("bottomright", legend = paste0(modelNames, ": ", round(auc, 3)),
      col = 1:4, lwd = 2)

```



Linear

```
linear_svm <- tune.svm(quality~.,
  data = wine[rowtrain,],
  kernel = "linear",
  cost = data.frame(cost = exp(seq(-20,-15,len=50))))
linear_svm$best.parameters
```

```
##          cost
## 1 2.061154e-09
```

Radial

```
radi_svm <- tune.svm(quality~.,
  data = wine[rowtrain,],
  kernel = "radial",
  cost = data.frame(cost = exp(seq(-20,-15,len=50))))
radi_svm$best.parameters
```

```
##          cost
## 1 2.061154e-09
```

only error rate can be obtained

```
linea.svm.pred = predict(linear_svm$best.model, newdata = wine[-rowtrain,], type = "prob")
radial.svm.pred = predict(radi_svm$best.model, newdata = wine[-rowtrain,], type = "prob")

error_linear=1-sum(linea.svm.pred==wine[-rowtrain,12])/nrow(wine[-rowtrain,])
error_radial=1-sum(radial.svm.pred==wine[-rowtrain,12])/nrow(wine[-rowtrain,])

tree.pred2 = predict(tree2, newdata = wine[-rowtrain,])
bag.pred2 = predict(bagging, newdata = wine[-rowtrain,])
rf.pred2 = predict(rf, newdata = wine[-rowtrain,])

error_tree=1-sum(tree.pred2 ==wine[-rowtrain,12])/nrow(wine[-rowtrain,])
error_bag=1-sum(bag.pred2==wine[-rowtrain,12])/nrow(wine[-rowtrain,])
error_rf=1-sum(rf.pred2==wine[-rowtrain,12])/nrow(wine[-rowtrain,])

data.frame(model=c("SVM_linear","SVM_radial","Tree","Bagging","Random_forest"),error_rate=c(error_linear,
```

model	error_rate
SVM_linear	0.3667436
SVM_radial	0.3667436
Tree	1.0000000
Bagging	0.1852194
Random_forest	0.1704388